

Branch: master ▾

CSE-255 / homework 7 / 1.CoverType.ipynb

Find file

Copy path

 kophy homework 7

22eef7c on Jun 3, 2017

1 contributor

729 lines (728 sloc) | 125 KB

```
In [1]: import findspark
findspark.init()

from pyspark import SparkContext
sc = SparkContext(master="local[4]")

from pyspark.mllib.linalg import Vectors
from pyspark.mllib.regression import LabeledPoint

from string import split,strip

from pyspark.mllib.tree import GradientBoostedTrees, GradientBoostedTreesModel
from pyspark.mllib.tree import RandomForest, RandomForestModel
from pyspark.mllib.util import MLUtils
```

## Cover Type

Classify geographical locations according to their predicted tree cover:

- **URL:** <http://archive.ics.uci.edu/ml/datasets/Covertype> (<http://archive.ics.uci.edu/ml/datasets/Covertype>)
- **Abstract:** Forest CoverType dataset
- **Data Set Description:** <http://archive.ics.uci.edu/ml/machine-learning-databases/covtype/covtype.info> (<http://archive.ics.uci.edu/ml/machine-learning-databases/covtype/covtype.info>)

```
In [2]: #define a dictionary of cover types
CoverTypes={1.0: 'Spruce/Fir',
             2.0: 'Lodgepole Pine',
             3.0: 'Ponderosa Pine',
             4.0: 'Cottonwood/Willow',
             5.0: 'Aspen',
             6.0: 'Douglas-fir',
             7.0: 'Krummholz' }

print 'Tree Cover Types:'
CoverTypes
```

Tree Cover Types:

```
Out[2]: {1.0: 'Spruce/Fir',
        2.0: 'Lodgepole Pine',
        3.0: 'Ponderosa Pine',
        4.0: 'Cottonwood/Willow',
        5.0: 'Aspen',
        6.0: 'Douglas-fir',
        7.0: 'Krummholz'}
```

```
In [3]: # creating a directory called covtype, download and decompress covtype.data.gz into it

from os.path import exists
if not exists('covtype'):
    print "creating directory covtype"
    !mkdir covtype
%cd covtype
if not exists('covtype.data'):
    if not exists('covtype.data.gz'):
        print 'downloading covtype.data.gz'
        !curl -O http://archive.ics.uci.edu/ml/machine-learning-databases/covtype/covtype.data.gz
    print 'decompressing covtype.data.gz'
    !gunzip -f covtype.data.gz
```

```
/home/kophy/Github/CSE255-DSE230/Classes/HW-7/covtype
total 73408
-rw-rw-r-- 1 kophy kophy 75169317 Jun  2 20:44 covtype.data
/home/kophy/Github/CSE255-DSE230/Classes/HW-7
```

```
In [5]: # Break up features that are made out of several binary features.
from string import split,strip
cols=[strip(a) for a in split(cols_txt,',')]
colDict={a:[a] for a in cols}
colDict['Soil_Type (40 binary columns)'] = ['ST_'+str(i) for i in range(40)]
colDict['Wilderness_Area (4 binarycolumns)'] = ['WA_'+str(i) for i in range(4)]
Columns=[]
for item in cols:
    Columns=Columns+colDict[item]
print Columns
```

```
In [6]: # Have a look at the first two lines of the data file
!head -2 covtype/covtype.data
```

```
In [7]: # Read the file into an RDD
        # If doing this on a real cluster, you need the file to be available on all nodes, ideally in HDFS
        .
        path='covtype/covtype.data'
        inputRDD=sc.textFile(path).cache()
        inputRDD.first()
```

```
In [8]: # Transform the text RDD into an RDD of LabeledPoints
Data=inputRDD.map(lambda line: [float(strip(x)) for x in line.split(',')]).map(lambda x: (x[-1], x[:-1]))
Data.first()
```

```
Out[8]: (5.0,
         [2596.0,
          51.0,
          3.0,
          258.0,
          0.0,
          510.0,
          221.0,
          232.0,
          148.0,
          6279.0,
          1.0,
          0.0,
          0.0,
          0.0,
          0.0,
          0.0])
```

[illegible]

```
In [9]: # count the number of examples of each type
total=Data.cache().count()
print 'total data size=',total
counts = Data.map(lambda (k,v): (k,1)).reduceByKey(lambda v1, v2: v1 + v2).collect()
counts.sort(key = lambda x:x[1],reverse = True)
print '                type (label):    percent of total'
print '-----'
print '\n'.join(['%20s (%3.1f):\t%4.2f'%(CoverTypes[a[0]],a[0],100.0*a[1]/float(total)) for a in c
ounts])
```

```
total data size= 581012
      type (label):    percent of total
-----
      Lodgepole Pine (2.0):    48.76
      Spruce/Fir (1.0):    36.46
      Ponderosa Pine (3.0):    6.15
      Krummholz (7.0):    3.53
      Douglas-fir (6.0):    2.99
      Aspen (5.0):    1.63
      Cottonwood/Willow (4.0):    0.47
```

## Making the problem binary

The implementation of `BoostedGradientTrees` in `MLlib` supports only binary problems. the `CovType` problem has 7 classes. To make the problem binary we choose the `Lodgepole Pine` (label = 2.0). We therefor transform the dataset to a new dataset where the label is 1.0 if the class is `Lodgepole Pine` and is 0.0 otherwise.

```
In [10]: Label=2.0

def getBinaryLabel(label):
    return 1.0 if label == Label else 0.0

Data=inputRDD.map(lambda line: [float(x) for x in line.split(',')])\
               .map(lambda V:LabeledPoint(getBinaryLabel(V[-1]), V[:-1]))
```

## Reducing data size

In order to see the effects of overfitting more clearly, we reduce the size of the data by a factor of 10

```
In [11]: Data1=Data.sample(False,0.1).cache()
(trainingData,testData)=Data1.randomSplit([0.7,0.3], seed=255)

print 'Sizes: Data1=%d, trainingData=%d, testData=%d'%(Data1.count(),trainingData.cache().count(),
testData.cache().count())

Sizes: Data1=57888, trainingData=40740, testData=17148

In [12]: counts=testData.map(lambda lp:(lp.label,1)).reduceByKey(lambda x,y:x+y).collect()
counts.sort(key=lambda x:x[1],reverse=True)
counts

Out[12]: [(0.0, 8764), (1.0, 8384)]
```

## Gradient Boosted Trees

- Following [this example](http://spark.apache.org/docs/latest/mllib-ensembles.html#gradient-boosted-trees-gbts) (<http://spark.apache.org/docs/latest/mllib-ensembles.html#gradient-boosted-trees-gbts>) from the mllib documentation
- [pyspark.mllib.tree.GradientBoostedTrees](http://spark.apache.org/docs/latest/api/python/pyspark.mllib.html#pyspark.mllib.tree.GradientBoostedTrees) [documentation](http://spark.apache.org/docs/latest/api/python/pyspark.mllib.html#pyspark.mllib.tree.GradientBoostedTrees)

### Main classes and methods

- GradientBoostedTrees is the class that implements the learning trainClassifier,
  - It's main method is trainClassifier(trainingData) which takes as input a training set and generates an instance of GradientBoostedTreesModel
  - The main parameter from train Classifier are:
    - **data** – Training dataset: RDD of LabeledPoint. Labels should take values {0, 1}.
    - categoricalFeaturesInfo – Map storing arity of categorical features. E.g., an entry (n -> k) indicates that feature n is categorical with k categories indexed from 0: {0, 1, ..., k-1}.
    - **loss** – Loss function used for minimization during gradient boosting. Supported: {"logLoss" (default), "leastSquaresError", "leastAbsoluteError"}.
    - **numIterations** – Number of iterations of boosting. (default: 100)
    - **learningRate** – Learning rate for shrinking the contribution of each estimator. The learning rate should be between in the interval (0, 1]. (default: 0.1)
    - **maxDepth** – Maximum depth of the tree. E.g., depth 0 means 1 leaf node; depth 1 means 1 internal node + 2 leaf nodes. (default: 3)
    - **maxBins** – maximum number of bins used for splitting features (default: 32) DecisionTree requires maxBins >= max categories
- GradientBoostedTreesModel represents the output of the boosting process: a linear combination of classification trees. The methods supported by this class are:
  - save(sc, path) : save the tree to a given filename, sc is the Spark Context.
  - load(sc,path) : The counterpart to save - load classifier from file.
  - predict(X) : predict on a single datapoint (the .features field of a LabeledPont) or an RDD of datapoints.
  - toDebugString() : print the classifier in a human readable format.

```
In [16]: from time import time
errors={}
for depth in [1, 3, 6, 10]:
    start=time()
    model=GradientBoostedTrees.trainClassifier(trainingData, categoricalFeaturesInfo={},
maxDepth=depth, numIterations=10)

    #print model.toDebugString()
    errors[depth]={}
    dataSets={'train':trainingData,'test':testData}
    for name in dataSets.keys(): # Calculate errors on train and test sets
        data=dataSets[name]
        Predicted=model.predict(data.map(lambda x: x.features))
        LabelsAndPredictions=data.map(lambda x: x.label).zip(Predicted)
        Err = LabelsAndPredictions.filter(lambda (v,p):v != p).count()/float(data.count())
        errors[depth][name] = Err
    print depth,errors[depth],int(time()-start),'seconds'
print errors

1 {'test': 0.26831117331467225, 'train': 0.2721158566519391} 10 seconds
```

```

3 {'test': 0.24708420807091205, 'train': 0.24700540009818361} 10 seconds
6 {'test': 0.2201422906461395, 'train': 0.20986745213549338} 13 seconds
10 {'test': 0.17168182878469793, 'train': 0.13912616593028965} 23 seconds
1: {'test': 0.26831117331467225, 'train': 0.2721158566519391}, 10: {'test': 0.17168182878469793,
'train': 0.13912616593028965}, 3: {'test': 0.24708420807091205, 'train': 0.24700540009818361}, 6:
{'test': 0.2201422906461395, 'train': 0.20986745213549338}

```

In [17]: B10 = errors

```

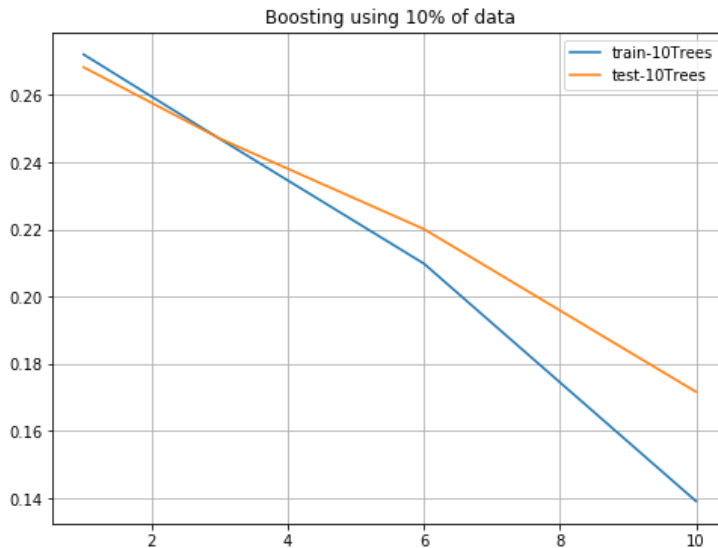
In [18]: # Plot Train/test accuracy vs Depth of trees graph
%pylab inline
from plot_utils import *
make_figure([B10],['10Trees'],Title='Boosting using 10% of data')

```

Populating the interactive namespace from numpy and matplotlib

/usr/local/lib/python2.7/dist-packages/IPython/core/magics/pylab.py:161: UserWarning: pylab import has clobbered these variables: ['e', 'split']

`%matplotlib` prevents importing \* from pylab and numpy  
 "\n`matplotlib` prevents importing \* from pylab and numpy"



## Random Forests

- Following [this example \(http://spark.apache.org/docs/latest/mllib-ensembles.html#classification\)](http://spark.apache.org/docs/latest/mllib-ensembles.html#classification) from the mllib documentation
- [pyspark.mllib.tree.RandomForest](http://spark.apache.org/docs/latest/api/python/pyspark.mllib.html#pyspark.mllib.tree.RandomForest) [documentation](http://spark.apache.org/docs/latest/api/python/pyspark.mllib.html#pyspark.mllib.tree.RandomForest)

**trainClassifier**(data, numClasses, categoricalFeaturesInfo, numTrees, featureSubsetStrategy='auto', impurity='gini', maxDepth=4, maxBins=32, seed=None)

Method to train a decision tree model for binary or multiclass classification.

### Parameters:

- **data** – Training dataset: RDD of LabeledPoint. Labels should take values {0, 1, ..., numClasses-1}.
- **numClasses** – number of classes for classification.
- **categoricalFeaturesInfo** – Map storing arity of categorical features. E.g., an entry (n -> k) indicates that feature n is categorical with k categories indexed from 0: {0, 1, ..., k-1}.
- **numTrees** – Number of trees in the random forest.
- **featureSubsetStrategy** – Number of features to consider for splits at each node. Supported: “auto” (default), “all”, “sqrt”, “log2”, “onethird”. If “auto” is set, this parameter is set based on numTrees: if numTrees == 1, set to “all”; if numTrees > 1 (forest) set to “sqrt”.
- **impurity** – Criterion used for information gain calculation. Supported values: “gini” (recommended) or “entropy”.
- **maxDepth** – Maximum depth of the tree. E.g., depth 0 means 1 leaf node; depth 1 means 1 internal node + 2 leaf nodes. (default: 4)
- **maxBins** – maximum number of bins used for splitting features (default: 32)
- **seed** – Random seed for bootstrapping and choosing feature subsets.

### Returns:

RandomForestModel that can be used for prediction

```

In [20]: from time import time
errors={}
for depth in [1, 2, 4, 10, 15, 20]:

```

```

for depth in [1,3,6,10,15,20]:
    start=time()
    model = RandomForest.trainClassifier(trainingData, numClasses=2, categoricalFeaturesInfo={},
                                       numTrees=10, featureSubsetStrategy="auto",
                                       impurity="gini", maxDepth=depth, maxBins=32)

    #print model.toDebugString()
    errors[depth]={}
    dataSets={'train':trainingData,'test':testData}
    for name in dataSets.keys(): # Calculate errors on train and test sets
        data=dataSets[name]
        Predicted=model.predict(data.map(lambda x: x.features))
        LabelsAndPredictions=data.map(lambda x: x.label).zip(Predicted)
        Err = LabelsAndPredictions.filter(lambda (v,p):v != p).count()/float(data.count())
        errors[depth][name]=Err
    print depth,errors[depth],int(time()-start),'seconds'
print errors

```

```

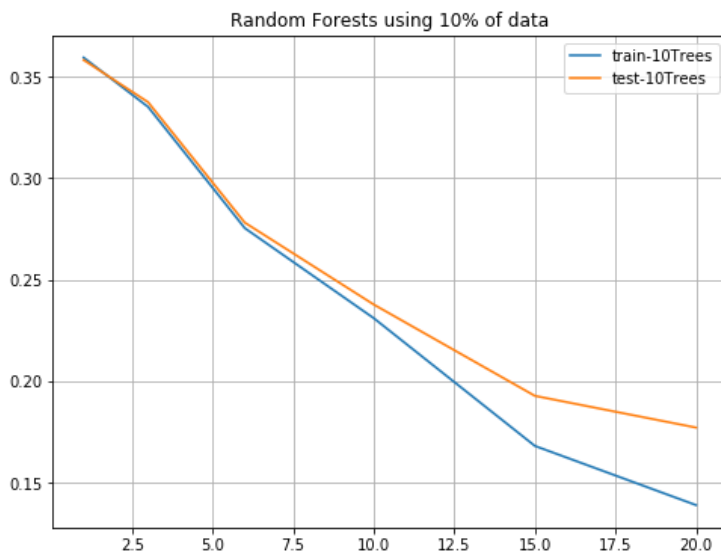
1 {'test': 0.35823419640774434, 'train': 0.35945017182130584} 2 seconds
3 {'test': 0.3374737578726382, 'train': 0.33512518409425623} 2 seconds
6 {'test': 0.27810823419640773, 'train': 0.2753804614629357} 3 seconds
10 {'test': 0.23775367389783064, 'train': 0.23102601865488465} 5 seconds
15 {'test': 0.19273384651271286, 'train': 0.16806578301423664} 8 seconds
20 {'test': 0.17710520177280148, 'train': 0.1389052528227786} 15 seconds
1: {'test': 0.35823419640774434, 'train': 0.35945017182130584}, 3: {'test': 0.3374737578726382, '
train': 0.33512518409425623}, 6: {'test': 0.27810823419640773, 'train': 0.2753804614629357}, 10: {
'test': 0.23775367389783064, 'train': 0.23102601865488465}, 15: {'test': 0.19273384651271286, 'tra
in': 0.16806578301423664}, 20: {'test': 0.17710520177280148, 'train': 0.1389052528227786}}

```

```

In [21]: RF_10trees = errors
# Plot Train/test accuracy vs Depth of trees graph
make_figure([RF_10trees],[ '10Trees'],Title='Random Forests using 10% of data')

```



Now plot B10 and RF\_10trees performance curves in the same graph

```

In [23]: make_figure([B10,RF_10trees],[ 'B_10Trees','RF_10Trees'],Title='Boosting and Random Forests using 10% of data')

```

